CS 83: Computer Vision

Image Filtering and Hough Transform

Amittai Siavava

2023-01-22

Abstract

This project explores image manipulation with convolutions and the Hough transform.

Contents

1. Theory Questions

Problem 1.

Show that if you use the line equation $\rho = x \cos \theta + y \sin \theta$, each image point x, y results in a sinusoid in (ρ, θ) Hough space. Relate the amplitude and the phase of the sinusoid to the point (x, y).

For an arbitrary point (x, y), its distance from the origin is $r = \sqrt{x^2 + y^2}$ and the equation of the line in (ρ, θ) space is $\rho = x \cos \theta + y \sin \theta$. We can derive the equivalent sinusoid:

$$\rho = x \cos \theta + y \sin \theta$$

$$\rho = r \left(\frac{x \cos \theta}{r} + \frac{y \sin \theta}{r} \right)$$

$$\rho = r \left(\sin \gamma \cos \theta + \cos \gamma \sin \theta \right) \qquad \text{(since } x, y, \text{ and } r \text{ form a right-angle triangle)}$$

$$\rho = r \sin(\gamma + \theta)$$

Therefore, the equivalent sinusoid has amplitude $r = \sqrt{x^2 + y^2}$ and a phase-shift $\gamma = \arctan \left(\frac{x}{y} \right)$.

Problem 2.

Why do we parametrize the line in terms of (ρ, θ) instead of the slope and the intercept, (m, c)?

Representing the line in terms of (ρ, θ) reduces the space of the parameters. In slope form $(y = mx + c)$, the slope m and the intercept c have a range of $[-\infty, \infty]$. The accumulator needed to store the votes for each possible line is therefore unbounded. In (ρ, θ) form, θ is an angle in the range $[0, 360]$ and ρ is bounded by the longest possible distance in the image, $\sqrt{W^2 + H^2}$, allowing us to more feasibly compute the accumulator array and find lines in the image.

Express the slope and the intercept in terms of (ρ, θ).

Recall that $\rho = x \cos \theta + y \sin \theta$. by re-arranging the equation, we get:

$$x \cos \theta + y \sin \theta = \rho \qquad\qquad y = mx + c$$

$$y \sin \theta = \rho - x \cos \theta \qquad\qquad \implies m = -\frac{\cos \theta}{\sin \theta}$$

$$y = \frac{\rho}{\sin \theta} - \frac{\cos \theta}{\sin \theta} x \qquad\qquad \implies c = \frac{\rho}{\sin \theta}$$

3

Problem 3.

Assuming that the image points (x, y) are in an image of width W and height H, that is $x \in [1, W]$ and $y \in [1, H]$, what is the maximum absolute value of ρ, and what is the range of θ?

The maximum absolute value of ρ is the longest distance that can fit in the image, $\sqrt{W^2 + H^2}$.

The range of θ is $[0, 360]$.

Problem 4.

For point $(10, 10)$ and points $(20, 20)$ and $(30, 30)$ in the image, plot the corresponding sinusoid waves in Hough space and visualize how their intersection point defines the line. What is (m, c) for this line?

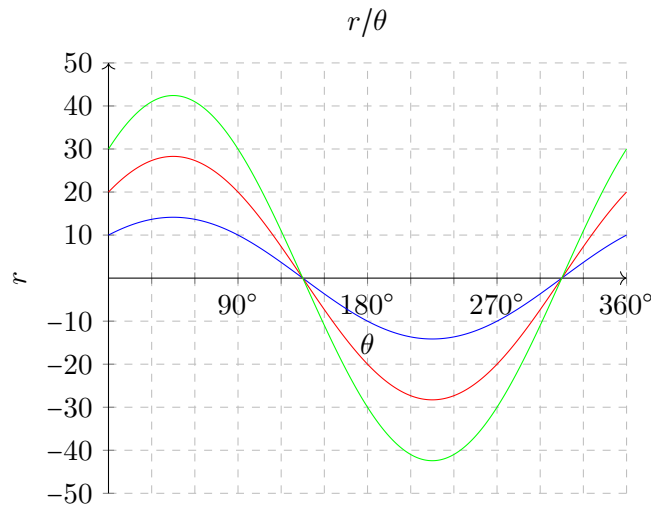| Point (x, y)	$r = \sqrt{x^2 + y^2}$	$\gamma = \arctan(x/y)$	Sinusoid
$(10, 10)$	$\sqrt{200}$	45	$\sqrt{200} \sin(\theta + 45)$
$(20, 20)$	$\sqrt{800}$	45	$\sqrt{800} \sin(\theta + 45)$
$(30, 30)$	$\sqrt{1800}$	45	$\sqrt{1800} \sin(\theta + 45)$

TABLE 1. Point and sinusoid values

FIGURE 1. Generated sinusoidal waves

The intersection points are at $r = 0$ and $\theta \in \{135, 315\}$.

$$\cos\theta = -\sin\theta$$

$$\rho = x\cos\theta + y\sin\theta = x - y = 0 \qquad\qquad \text{(since } x = y\text{)}$$

$$m = -\frac{\cos\theta}{\sin\theta} = 1$$

$$c = \frac{\rho}{\sin\theta} = 0$$

2. Experiments and Results

2.1. Design Challenges.

(i) I was not sure if this is a standard change, but `scipy.signal.gaussian` seems to have been migrated to `scipy.signal.windows.gaussian` in Python 3.11 (see here).

(ii) I spent a few hours debugging my hough transform, only to realize that writing into and reading from an array concurrently was resulting in inaccuracies. Fixing this issue made everything work.

(iii) I also had a bug in my gaussian kernel — I was resizing the kernel but not normalizing it back to sum up to 1, which resulted in a lot of spurious and inaccurate lines. This was a definitive fix, but debugging it took a while.

2.2. Experiments. I played around with some of the parameters to see how the results are affected.

(i) **gaussian kernel variance (σ):** Increasing the value of σ improves the image detection up to a point, then flattens out the line once σ goes beyond the optimal value. I found the best results with $\sigma = 2$ to work well for most of my images, although $\sigma = 3$ seemed to work better on one image.

(ii) **threshold:** Raising the threshold prunes erratic lines in the image, but if the threshold is too high, accurate lines are pruned too (see image). I found the best results with a threshold of 0.3.

FIGURE 2. Threshold = 0.5

(iii) **theta resolution:** Increasing the resolution of θ increases the accuracy of the lines found in the image, but it also increases the computation time. I found the best results with a resolution of $\frac{\pi}{180}$, but I settled on $\frac{pi}{90}$ because the latter takes significantly less time to compute and the difference in accuracy is not too extreme.

(iv) **rho resolution:** I also found that increasing the resolution of ρ increases the accuracy of the lines found also increases the computation time.

(v) **number of lines:** This yielded varying results depending on the image. While a specific number of lines would cause the algorithm to ignore correct lines in one image, the same might cause less accurate lines to be admitted in another image. This clearly needs to be tuned specific to every image

2.3. **Results.** These were the results I got for image 01.


FIGURE 3. Original image

FIGURE 4. Edge Detection

FIGURE 5. Thresholding

FIGURE 6. Hough Transform

FIGURE 7. Line Detection